



Kubernetes on AWS: Cost optimization and reducing operational effort

Eryan Ariobowo
28 July 2020

© 2020, Amazon Web Services, Inc. or its Affiliates.



AWS Container Day at KubeCon

August 17, 2020 | 8:00AM - 4:00PM PDT

Register now

<https://bit.ly/2BC0461>

© 2020, Amazon Web Services, Inc. or its Affiliates.



Agenda

- Containers and Compute options
- Compute Cost and Performance
- Maximize Price/Performance
- Cost Optimization and Autoscaling
- Spot with Amazon EKS

Containers and compute options



Kubernetes on AWS and compute

Running K8s Options

Deployment, scheduling,
scaling, and management of
containerized applications



**Self-managed
Kubernetes on EC2**



**Amazon Elastic
Kubernetes Service
(Amazon EKS)**



**3rd Party-managed
Kubernetes**

Compute Engine

Where the containers run



**Amazon EC2
Self-Managed
Worker Nodes**



**AWS Fargate
for EKS**



**Amazon EC2
Managed Worker Nodes
on EKS**

Amazon Elastic Kubernetes Service



Upstream Kubernetes

Use standard Kubernetes APIs.
Works with community tools.

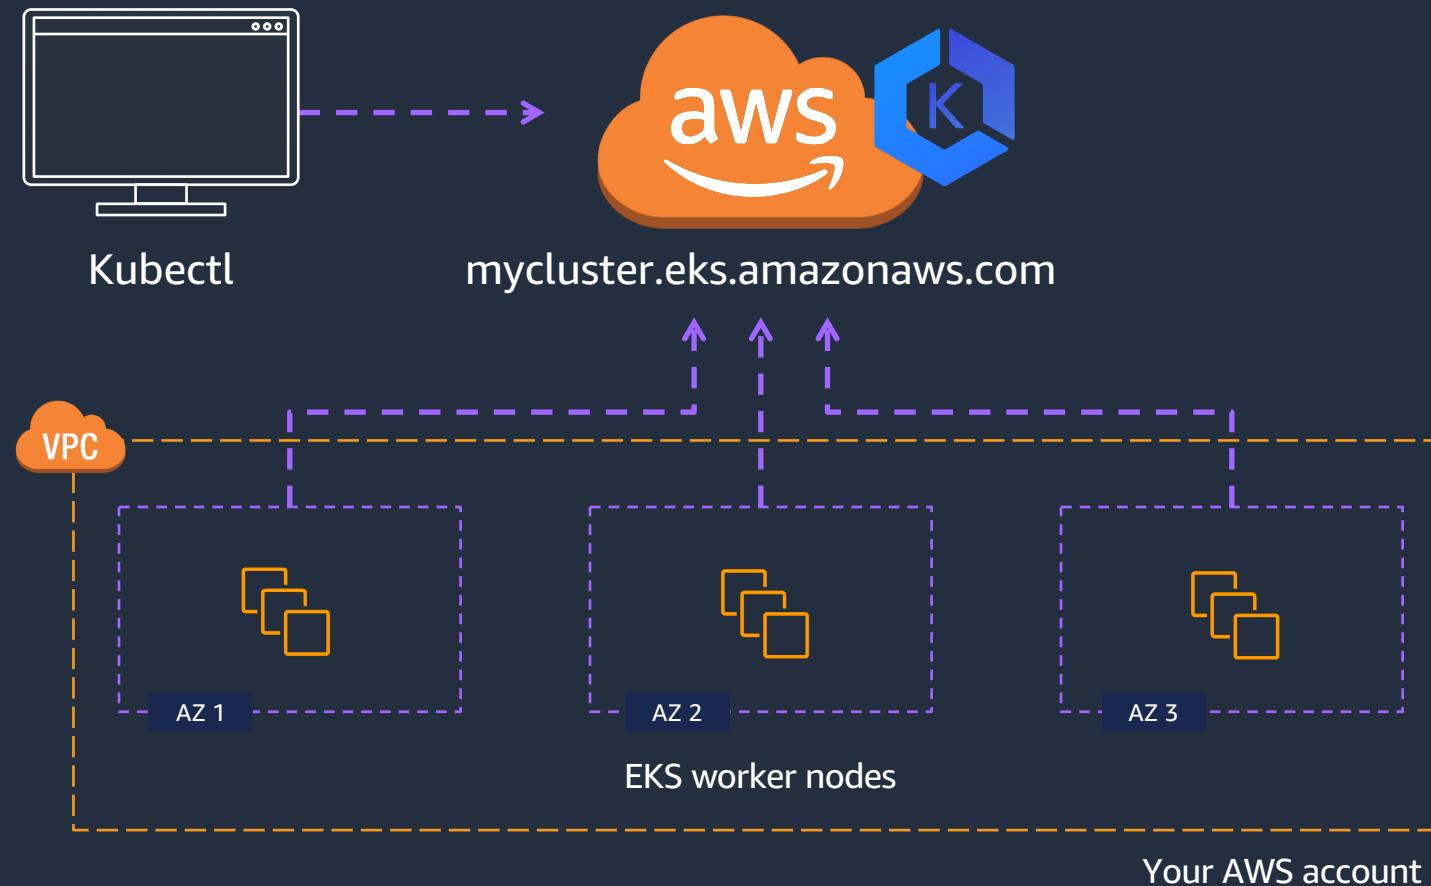
Highly available

Built for production workloads, all clusters are highly available. Backed by a 99.9% SLA.

Integrated

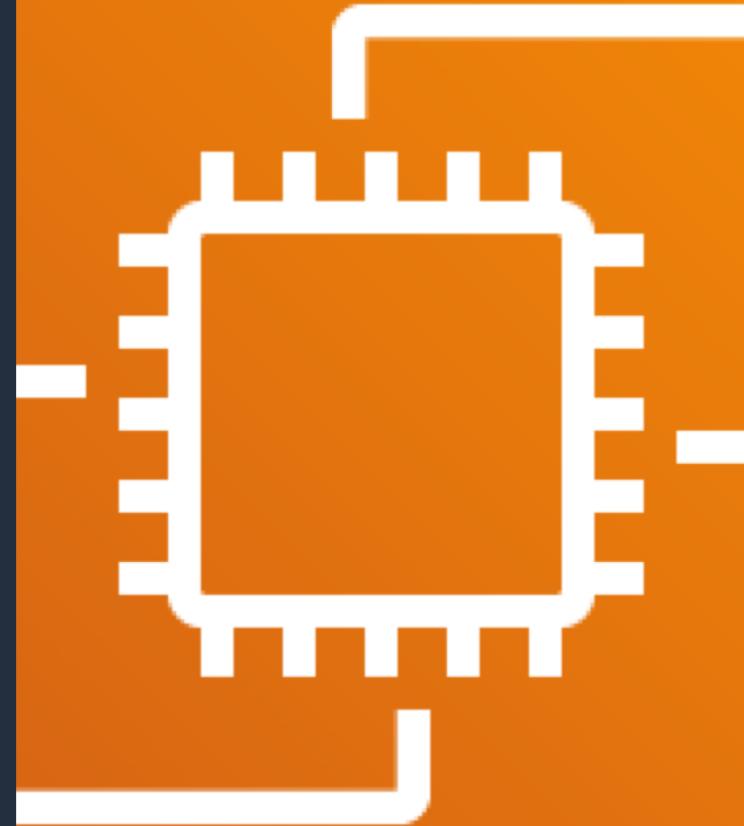
with the AWS ecosystem: VPC Networking, Elastic Load Balancing, IAM Permissions, CloudWatch and more

Amazon Elastic Kubernetes Service

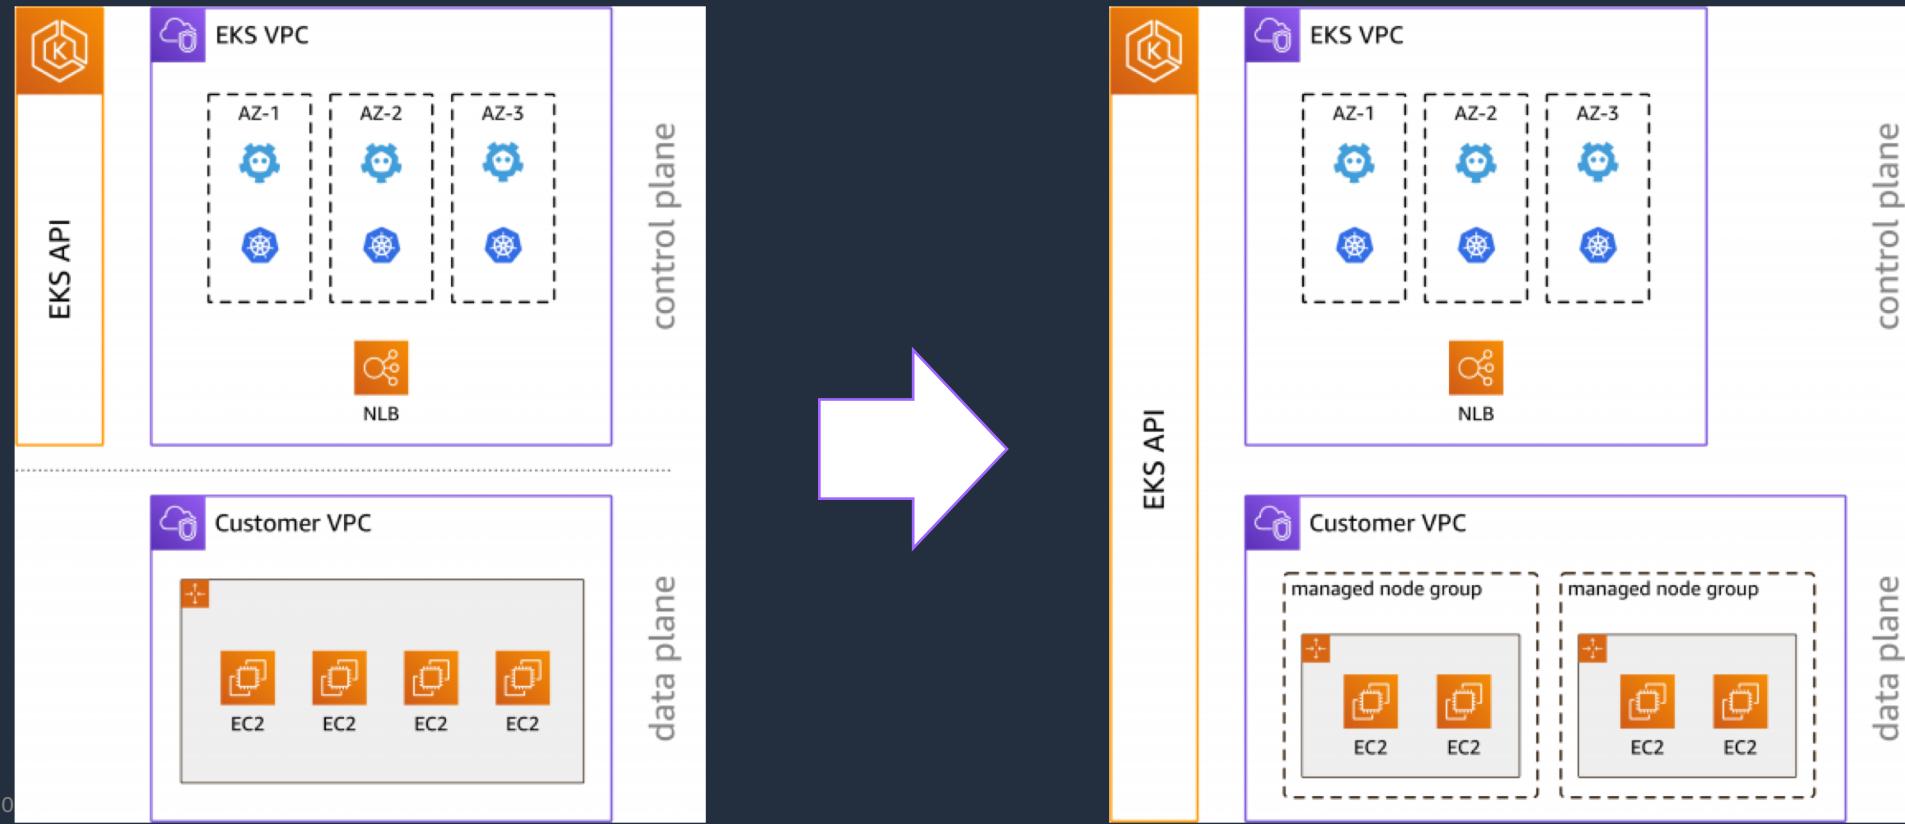


EC2 Mode – Customer Responsibilities

- Instance type and quantity to choose?
- Which OS to choose?
- Hardening the OS (e.g. against CIS benchmark) & Patching of the OS, Docker, kubelet etc.

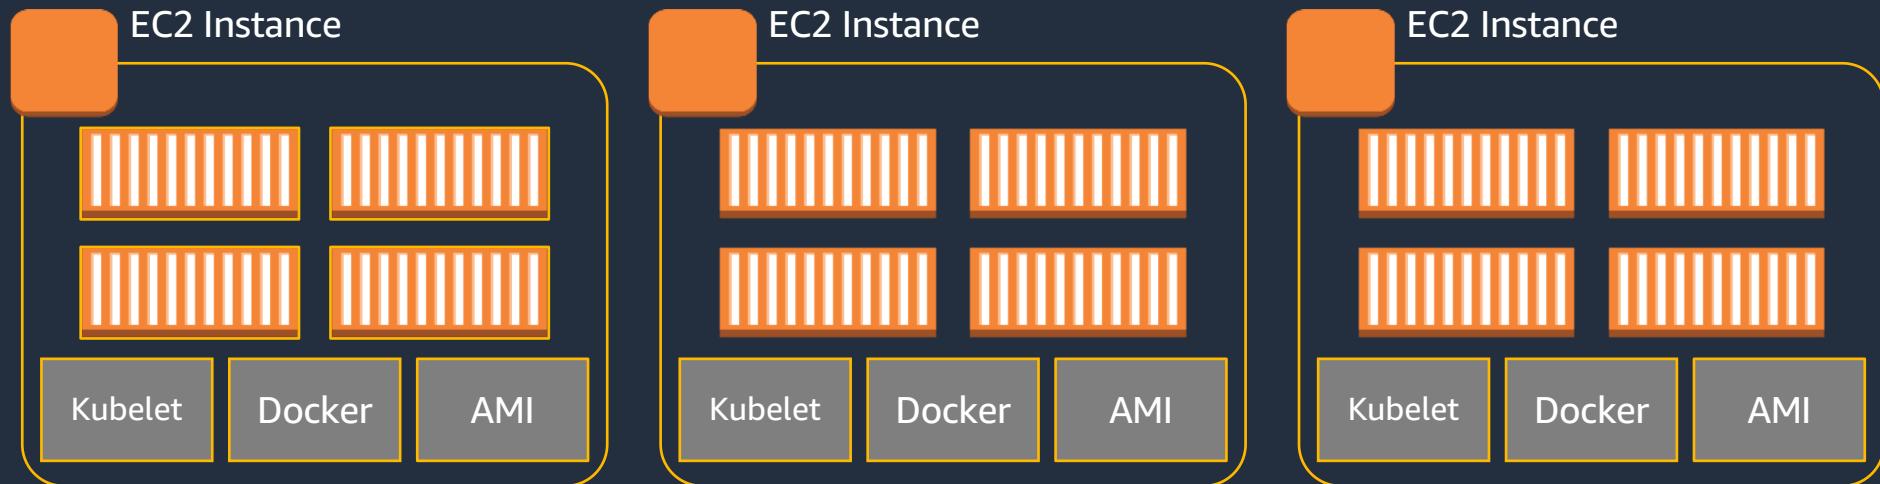


EKS managed worker nodes





Kubernetes Control Plane

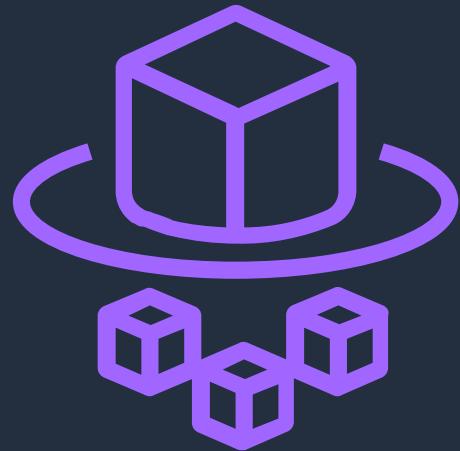




Kubernetes Control Plane



AWS Fargate



Managed by AWS

No EC2 Instances to provision, scale or manage

Elastic

Scale up & down seamlessly. Pay only for what you use

Integrated

with the AWS ecosystem: VPC Networking, Elastic Load Balancing, IAM Permissions, CloudWatch and more. Run Kubernetes pods or ECS tasks.

Security Benefits Of Fargate

We do more, you do less:

- Patching of the OS, Docker, ECS Agent, etc.
- Task isolation (via separate Clusters)
- Requires awsvpc network mode so ENI/SG per Task
- No runtime access for users (ssh or interactive Docker)



Cost Optimization



Cost optimization is about trading off cost and performance:
How can you achieve your performance goal at minimum cost?

Amazon EC2 purchase options

On-Demand

Pay-for-compute capacity **by the second** with no long-term commitments



Spiky workloads,
to define needs

Savings Plans & Reserved Instances

Make a commitment and receive a **significant discount** off compute



Committed &
steady-state usage

Spot Instances

Spare Amazon EC2 capacity at **savings of up to 90%** off On-Demand prices

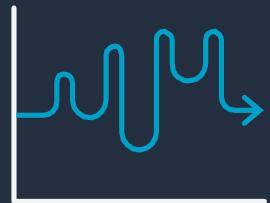


Fault-tolerant, flexible,
stateless workloads

Fargate Purchase Options

Fargate

Pay for containers **per-second** with no long-term commitment



Capacity needs can change rapidly

Compute Savings Plan

Make a 1 or 3-year commitment and receive a **significant discount**



Baseline compute needs known in advance

New

Fargate Spot

Spare capacity with **savings up to 70%** off Fargate standard pricing



Fault-tolerant, flexible workloads

Containers + Spot Instances



- Containers are often stateless, fault-tolerant, and a great fit for Spot Instances
- Deploy containerized workloads and easily manage clusters at any scale at a fraction of the cost with Spot Instances
- Spot instances can be used with Kubernetes to run any containerized workload



Skyscanner is a travel fare aggregator website and travel metasearch engine based in Edinburgh, Scotland

"We are currently tracking
74% saving over all regions."

- **Paul Gillespie,**
Principal Architect/Tribe Lead



Amazon EC2 Spot and Fargate Spot



Amazon EC2 Spot

Spare EC2 Capacity

Save up to **90%** over On-Demand

Can be *reclaimed* by Amazon EC2
(with two minute warning)

Follow EC2 Spot best practices –
instance flexibility and use
capacity optimized allocation
strategy for ASGs.



AWS Fargate Spot

Spare compute Capacity

Save up to **70%** over standard
Fargate

Can be *reclaimed*
(with two minute warning)

Automatic diversification

The Simple Rules of Spot



Spot infrastructure

Is same as On-Demand and RIs

70 - 90% off



Spot pricing

Smooth, infrequent changes
no spikes, more predictable



Interruptions

Only happen when OD needs capacity (**no bidding**)



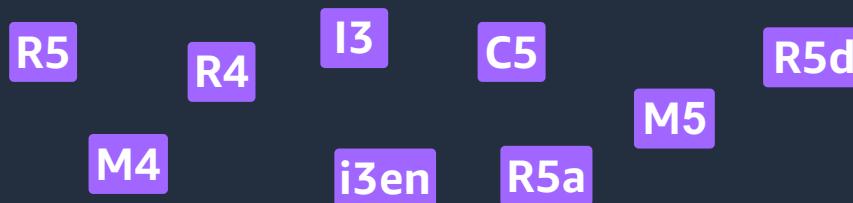
Diversify

Choose different instance types,
size and AZ in a single fleet

EC2 Spot instance pools explained

C4	1a	1b	1c	On Demand
8XL	\$0.50	\$0.27	\$0.29	\$1.76
4XL	\$0.21	\$0.30	\$0.16	\$0.88
2XL	\$0.08	\$0.07	\$0.08	\$0.44
XL	\$0.04	\$0.05	\$0.04	\$0.22
L	\$0.01	\$0.01	\$0.04	\$0.11

- Each instance family
- Each instance size
- Each Availability Zone
- In every region
- Is a separate Spot pool



Fargate Spot Pricing

Fargate		Fargate Spot	
per vCPU per hour	\$0.04048	per vCPU per hour	\$0. 012144
per GB per hour	\$0.004445	per GB per hour	\$0. 0013335

Supported Configurations	
CPU	Memory Values
0.25 vCPU	0.5GB, 1GB, and 2GB
0.5 vCPU	Min. 1GB and Max. 4GB, in 1GB increments
1 vCPU	Min. 2GB and Max. 8GB, in 1GB increments
2 vCPU	Min. 4GB and Max. 16GB, in 1GB increments
4 vCPU	Min. 8GB and Max. 30GB, in 1GB increments

What about interruptions?

Minimal interruptions

Over **95%** of the instances were not interrupted in the last 3 months



Leverage **Spot Instance Advisor** to select the suitable instance types to your workload with the lowest interruption rates

Catch the interruption notifications by polling EC2 metadata or using Cloudwatch events and **automate** response to interruptions:

- ELB connection draining and graceful shutdown
- Checkpointing
- **Cordon your node and drain containers**

Compute Cost and Performance

© 2020, Amazon Web Services, Inc. or its Affiliates.



The Components of Compute Cost

$$\text{(total compute cost)} = \text{(# of compute units)} \times \text{(time in use)} \times \text{(cost per unit)}$$

A diagram illustrating the formula for compute cost. On the left, a white outline of a money bag with a dollar sign inside is followed by an equals sign. To the right of the equals sign is a multiplication sign. The first factor is a yellow icon of two interlocking computer chips. The second factor is a yellow icon of a clock with a timeline above it. The third factor is a yellow icon of a dollar bill with a blue bracket underneath it. Below each icon is its corresponding label in white text.

What are Compute Units?

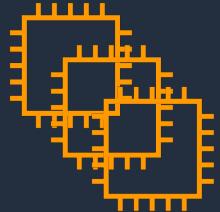


The Components of Compute Performance



=

(total compute performance)



X



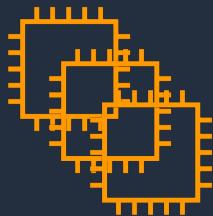
(# of compute units)

(performance per unit)

Minimize Cost While Meeting Performance Goals



{ (1) Choose compute units with the best *price/performance* for your workload.



{ (2) Run the *right number* for only *as long as you need them* to meet performance goals.

Cost Optimization and Autoscaling

© 2020, Amazon Web Services, Inc. or its Affiliates.



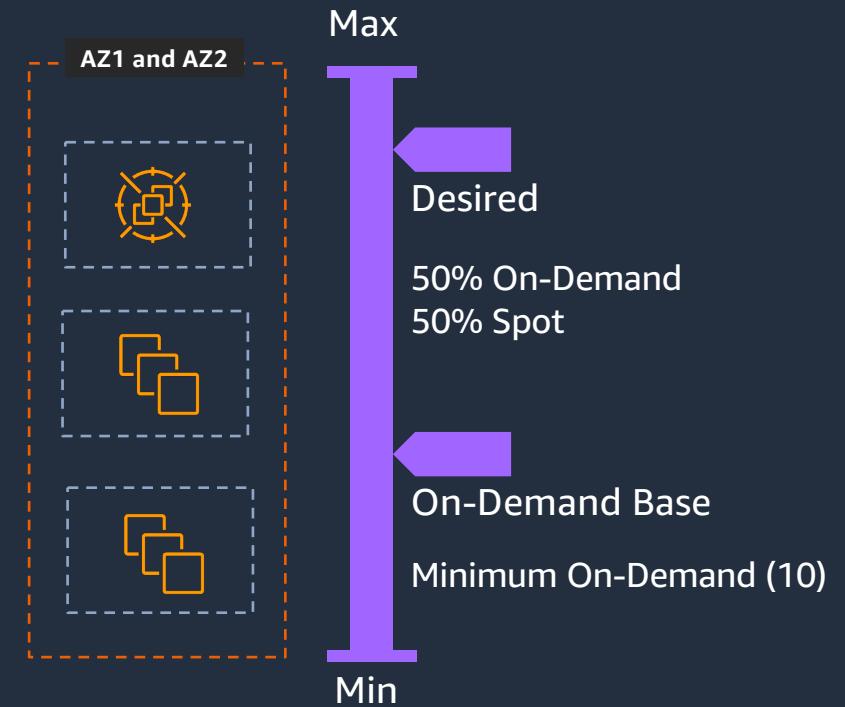
Amazon EC2 Auto Scaling groups



- a Support for **mixed instance types & Purchase Options**
- b Replacement of unhealthy & interrupted instances
- c Balanced capacity across availability zones
- d Elastic Load Balancer integration
- e Lifecycle hooks, termination policies, instance termination protection, scaling processes...

API Parameters

```
"MixedInstancesPolicy": {  
    "LaunchTemplate": {  
        "LaunchTemplateSpecification": {  
            "LaunchTemplateName": "MyLaunchTemplate"  
        },  
        "Overrides": [  
            { "InstanceType": "c5.large" },  
            { "InstanceType": "c4.large" }  
        ]  
    },  
    "InstancesDistribution": {  
        "OnDemandAllocationStrategy": "prioritized",  
        "OnDemandBaseCapacity": 10,  
        "OnDemandPercentageAboveBaseCapacity": 50,  
        "SpotAllocationStrategy": "capacity-optimized"  
    }  
}
```



ASG SpotAllocationStrategy: lowest-price (across N pools)

Desired capacity: 12

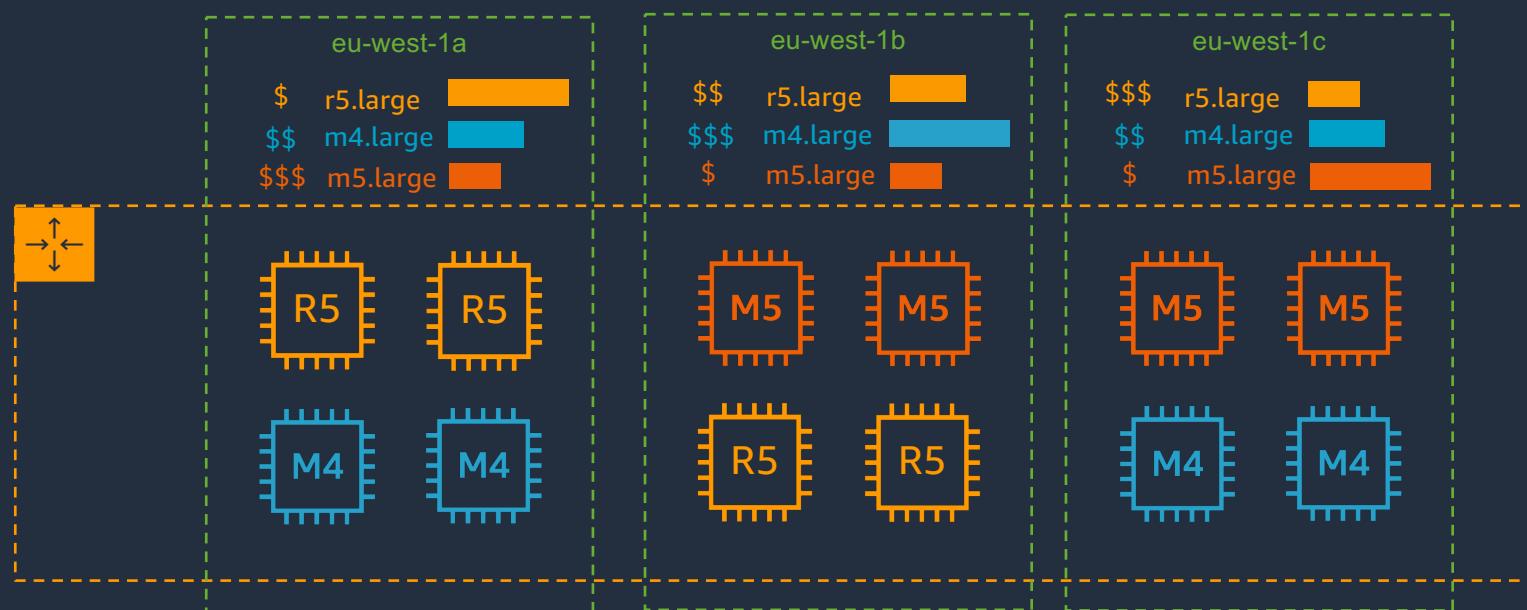
OnDemandBaseCapacity: 0

OnDemandPercentageAboveCapacity: 0

overrides: ["r5.large", "m4.large", "m5.large"]

SpotAllocationStrategy: lowest-price

SpotInstancePools: 2



ASG SpotAllocationStrategy: capacity-optimized

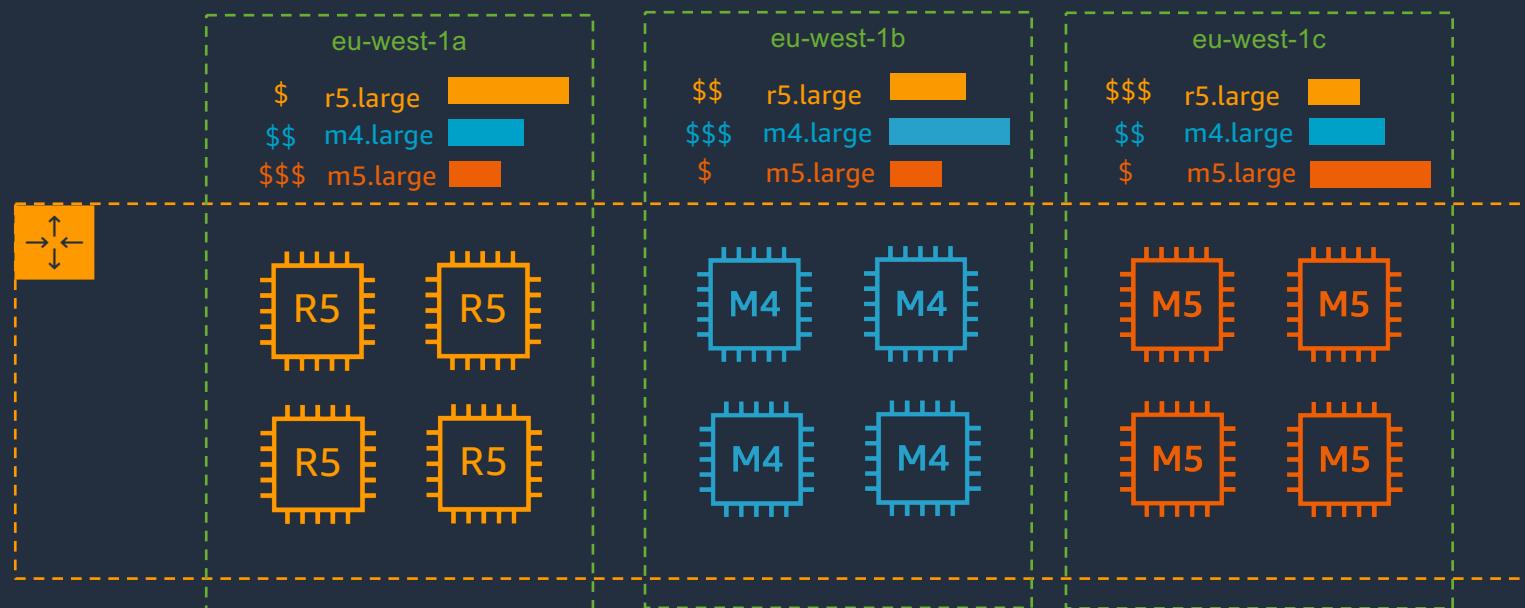
Desired capacity: 12

OnDemandBaseCapacity: 0

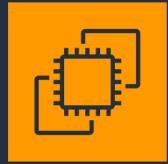
onDemandPercentageAboveCapacity: 0

overrides: ["r5.large", "m4.large", "m5.large"]

SpotAllocationStrategy: capacity-optimized



(1) Maximizing Price/Performance of Compute Units



EC2 Instances

Size task/pod vCPU and Memory
Choose Purchase Type
Choose Instance Type(s)



Fargate Tasks/Pods

Size task/pod vCPU and Memory
Choose Purchase Type

(1) Maximizing price/performance with 100% Spot



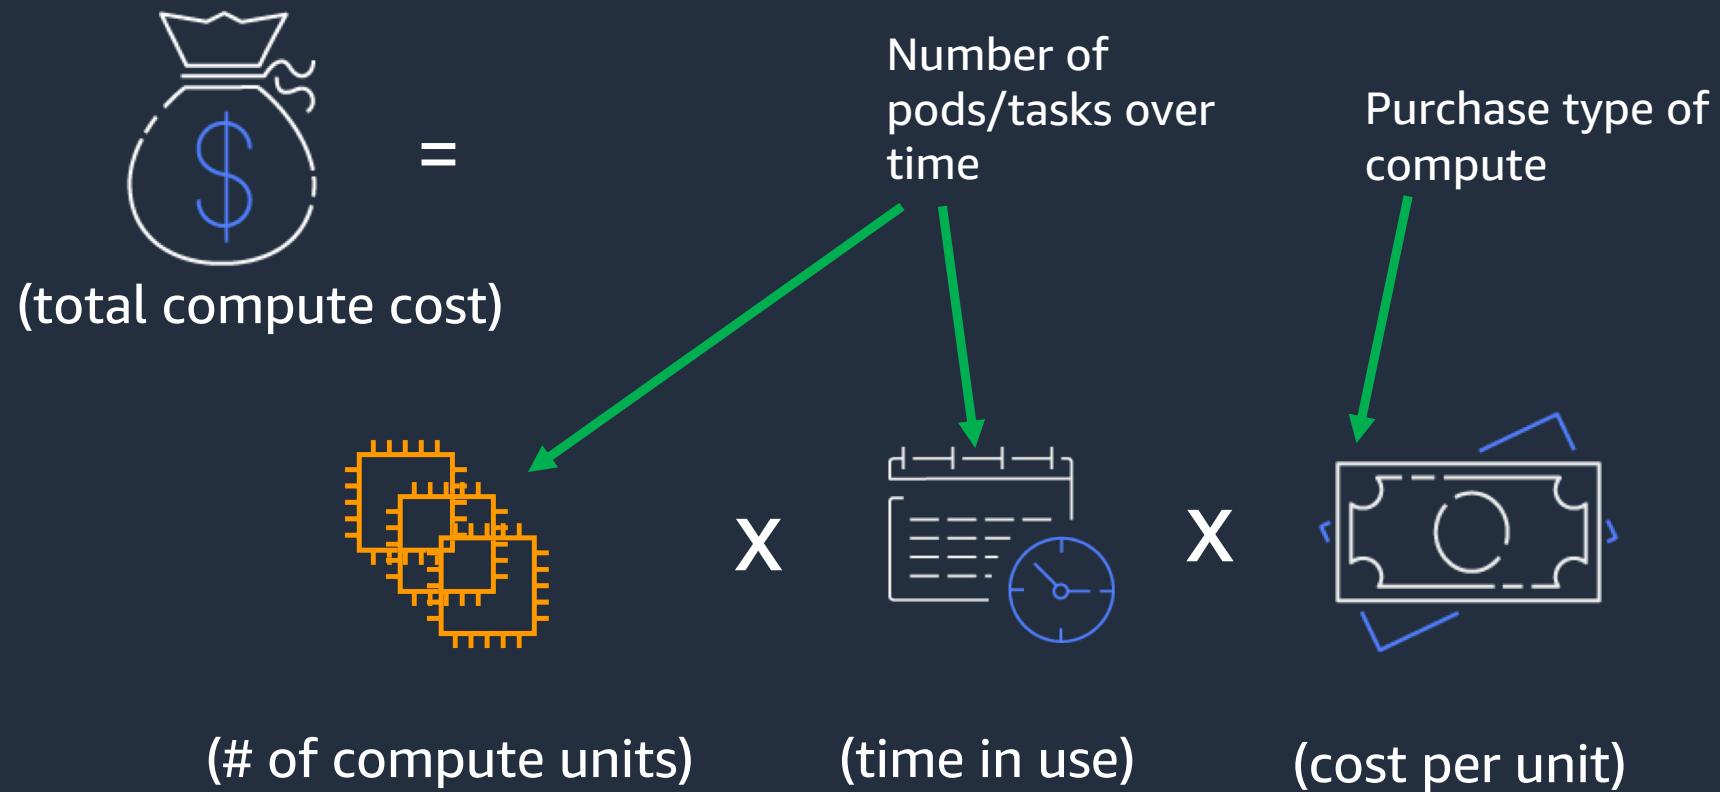
Good fit for 100% Spot:

- ETL
- Batch processing
- Dev/test environments

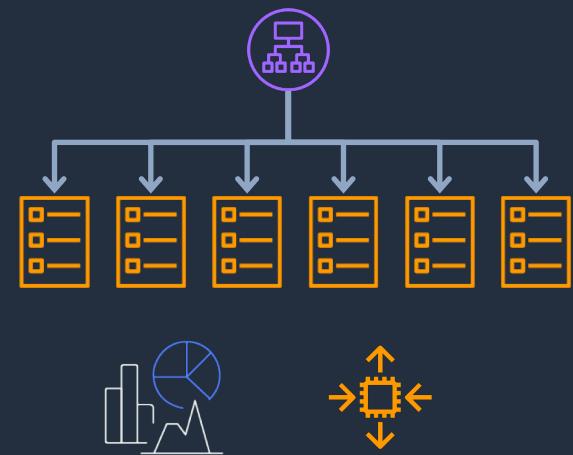
(2) Run the *right number* for only as long as needed



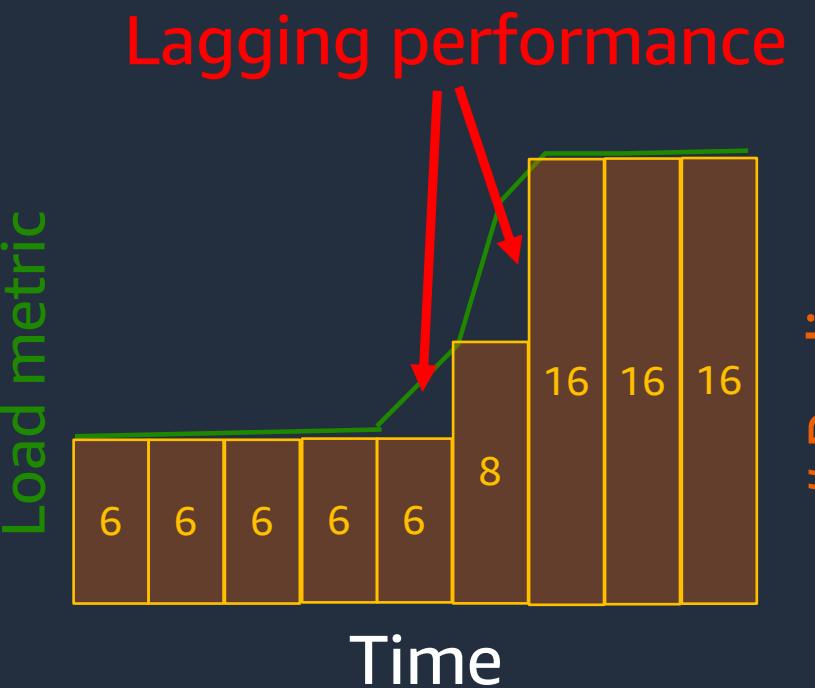
The Components of Compute Cost (revisited)



(2) Run the *right number* for only as long as needed

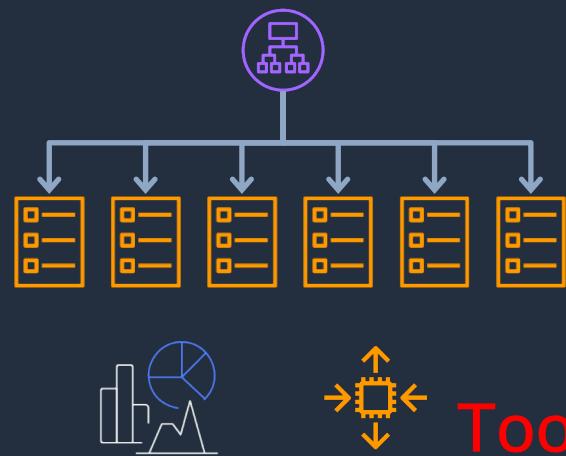


© 2020, Amazon Web Services, Inc. or its Affiliates.

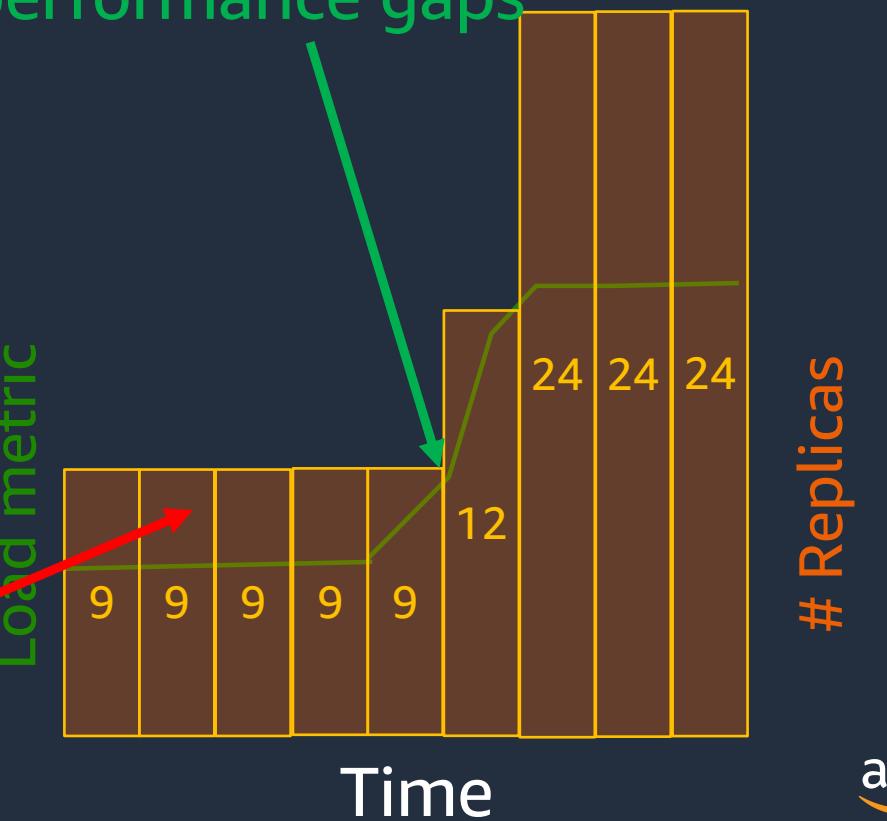


(2) Run the *right number* for only as long as needed

Overprovision by 50%:
Reduce metric target value by 1/3



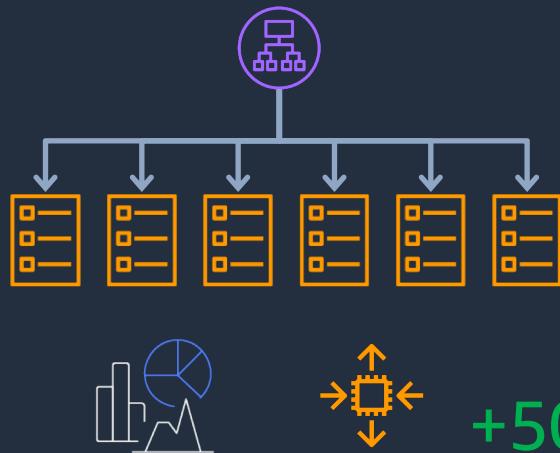
No performance gaps



(2) Run the *right number* for only as long as needed

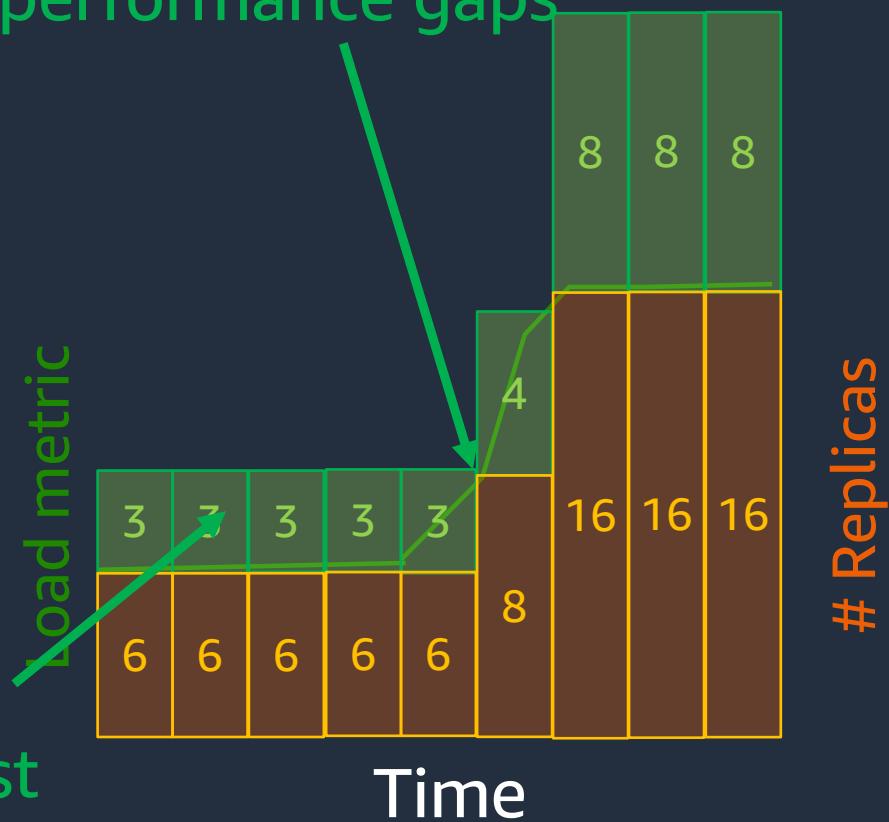
Overprovision by 50%:
Reduce metric target value by 1/3

Run 2/3 On-Demand, 1/3 on Spot



+50% capacity
for +5-10% cost

No performance gaps



Amazon EKS & EC2 Spot Instances

© 2020, Amazon Web Services, Inc. or its Affiliates.



Using EC2 Spot instances with Amazon EKS

1. Acquiring Capacity
2. Scaling Mechanisms: Horizontal Pod AutoScaler (HPA), Cluster Autoscaler (CA)
3. Interruption handling with a DaemonSet
4. Scheduling: Node Affinity, Taints and Tolerations

eksctl – Adding a diversified node group

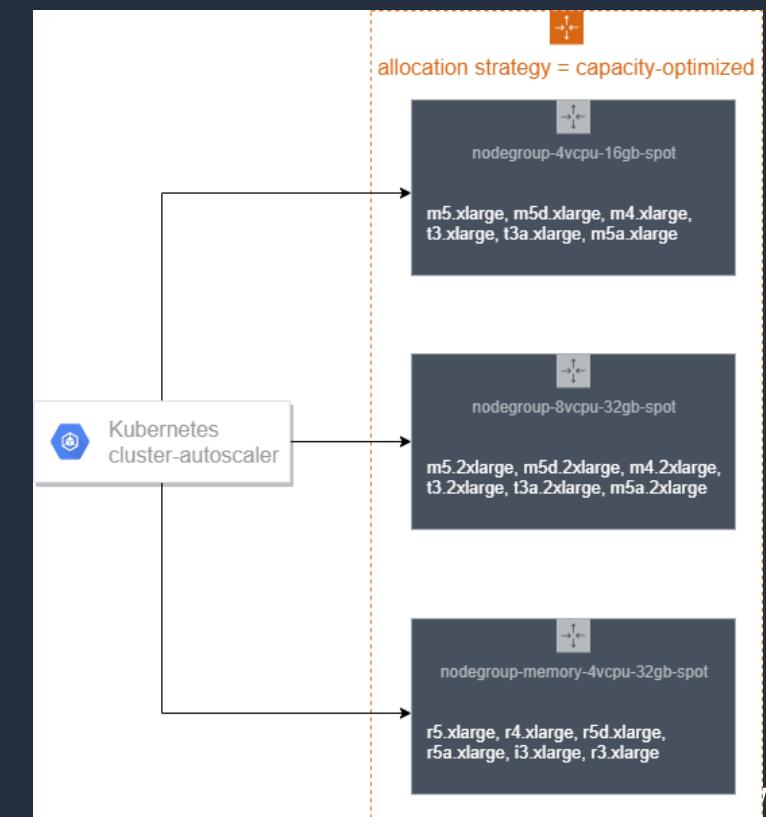
```
apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig
metadata:
  name: test-cluster
  region: us-west-2
nodeGroups:
  - name: dev-4vcpu-16gb-spot
    availabilityZones: ["us-west-2a", "us-west-2b", "us-west-2c"] 
    minSize: 1
    maxSize: 100
    instancesDistribution:
      instanceTypes: ["m5.xlarge", "m5d.xlarge", "m4.xlarge", "t3.xlarge", "t2.xlarge"] 
      onDemandBaseCapacity: 0
      onDemandPercentageAboveBaseCapacity: 0
      spotInstancePools: 4
    labels:
      lifecycle: Ec2Spot 
      environment: dev
      costcenter: engineering
      project: default
    taints:
      spotInstance: "true:PreferNoSchedule"
```

Auto Scaling the App & Cluster

CA Nodegroups are still expected to be homogeneous. Implement diversification !

command:

- ./cluster-autoscaler
- --v=4
- --stderrthreshold=info
- --cloud-provider=aws
- --skip-nodes-with-local-storage=false
- --nodes=0:100:nodegroup-dev-4vcpu-16gb-spot
- --nodes=0:100:nodegroup-dev-8vcpu-32gb-spot
- --nodes=0:100:nodegroup-dev-4vcpu-16gb-ondemand
- --nodes=0:100:nodegroup-dev-8vcpu-32gb-ondemand
- --expander=priority



AWS Interruptions Handler DaemonSet

<https://github.com/aws/aws-node-termination-handler>

```
log.Println("Kubernetes Spot Node Termination Handler has started successfully!")
waitForTermination(nthConfig)
drainer.Drain(nthConfig)
log.Printf("Node %q successfully drained.\n", nthConfig.NodeName)
```

<https://github.com/aws/eks-charts/tree/master/stable/aws-node-termination-handler>

```
helm repo add eks https://aws.github.io/eks-charts
helm install --name aws-node-termination-handler \
  --namespace kube-system eks/aws-node-termination-handler \
  --set nodeSelector.lifecycle=Ec2Spot
```

Taints, Toleration & Affinity

```
requiredDuringSchedulingIgnoredDuringExecution:  
  nodeSelectorTerms:  
    - matchExpressions:  
        - key: environment  
          operator: In  
          values:  
            - dev
```

Multi-tenant cluster. Group affinity

```
affinity:  
  nodeAffinity:  
    preferredDuringSchedulingIgnoredDuringExecution:  
      - weight: 1  
        preference:  
          matchExpressions:  
            - key: lifecycle  
              operator: In  
              values:  
                - OnDemand
```

Life-cycle affinity & Toleration

```
tolerations:  
  - key: "spotInstance"  
    operator: "Equal"  
    value: "true"  
    effect: "PreferNoSchedule"
```

Main Takeaways

- Use Fargate to reduce operational effort
- Get discount, use RI & Saving Plans
- Right Sizing
- Use only when you need it
- Cluster Auto-scaling
- Use Spot Instances

Thank you!

© 2020, Amazon Web Services, Inc. or its Affiliates.

